

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

In [2]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report

In [30]: df = pd.read_csv("bank.csv")
df.head()
Out[30]:
   (1162, 17)

In [31]: df.head()

Out[31]:
   age  job  marital  education  default  balance  housing  loan  contact  day  month  duration  campaign  pdays  previous  poutcome  deposit
0  59    0      1          1          0    2343      1      0      2      5      8    1042      1      -1      0      3      1
1  56    0      1          1          0     45      0      0      2      5      8    1467      1      -1      0      3      1
2  41    9      1          1          0   1270      1      0      2      5      8    1389      1      -1      0      3      1
3  55    7      1          1          0   2476      1      0      2      5      8    579      1      -1      0      3      1
4  54    0      1          2          0    184      0      0      2      5      8     673      2      -1      0      3      1

In [5]: df.deposit.value_counts()
Out[5]:
0      5873
1       5289
Name: deposit, dtype: int64

In [6]: x = df.iloc[:, :-1]
y = df.iloc[:, -1]

In [7]: xtrain, xtest, ytrain, ytest = train_test_split(x,y, test_size=0.3, random_state=1)

In [8]: def mymodel(model):
    model.fit(xtrain,ytrain)
    ypred=model.predict(xtest)
    print(classification_report(ytest,ypred))
    return model

In [9]: dt=DecisionTreeClassifier()

In [10]: mymodel(dt)

precision    recall  f1-score   support

   0.00      0.79      0.80      1760
   1.00      0.78      0.77      1589

 accuracy      0.79      3349
 macro avg      0.78      0.78      3349
 weighted avg      0.78      0.79      3349

DecisionTreeClassifier()

In [11]: dt.score(xtrain,ytrain)
Out[11]:
1.0

In [12]: dt.score(xtest,ytest)
Out[12]:
0.78501845088086

In [13]: dt.feature_importances_
array([0.08045239, 0.035509133, 0.01761428, 0.01205608, 0.00045636,
       0.00208469, 0.04261978, 0.00609448, 0.00214084, 0.07658374,
       0.00705057, 0.035031815, 0.02069599, 0.00673951, 0.01318065,
       0.03696931])
```

hyperparameter tuning

```
In [14]: dt1=DecisionTreeClassifier(max_depth=10)
mymodel(dt1)

precision    recall  f1-score   support

   0.00      0.83      0.83      1760
   1.00      0.81      0.81      1589

 accuracy      0.82      3349
 macro avg      0.82      0.82      3349
 weighted avg      0.82      0.82      3349

DecisionTreeClassifier(max_depth=10)

In [15]: for i in range(1,50):
    dt2=DecisionTreeClassifier(max_depth=i)
    dt2.fit(xtrain,ytrain)
    ypred=dt2.predict(xtest)
    print(f'{i}: {accuracy_score(ytest,ypred)}')

1: 0.7121528814571514
2: 0.7121528814571514
3: 0.7712750074640415
4: 0.78501845088086
5: 0.7993430874888027
6: 0.8006374738727978
7: 0.8147252172807417
8: 0.8181546739367274
9: 0.8118841445207524
10: 0.810647650917214
11: 0.8166619900567334
12: 0.8139743206927441
13: 0.80748051955807704
14: 0.8068360794687967
15: 0.798745934296805
16: 0.7884893982626421
17: 0.784269453568923
18: 0.78501845088086
19: 0.783607521648253
20: 0.787877954148408
21: 0.7868020304568528
22: 0.7852950134368408
23: 0.7868020304568528
24: 0.785607640728575
25: 0.7851988032248433
26: 0.7855031003034507
27: 0.7903851896688384
28: 0.7855936100328457
29: 0.7871006270529516
30: 0.7882960134368408
31: 0.7873992236488504
32: 0.7871006270529516
33: 0.7832188713048671
34: 0.7838108644968847
35: 0.7852040372640551
36: 0.78501845088086
37: 0.7855936100328457
38: 0.785607640728575
39: 0.7918781725888325
40: 0.787996416840848
41: 0.788692206268444
42: 0.785607640728575
43: 0.7936697521648253
44: 0.7852188713048671
45: 0.7891988032248433
46: 0.7853909474768588
47: 0.7842146810920625
48: 0.7817258883248731
49: 0.7921767691848313

In [16]: dt3=DecisionTreeClassifier(max_depth=11)
mymodel(dt3)

precision    recall  f1-score   support

   0.00      0.83      0.82      1760
   1.00      0.81      0.81      1589

 accuracy      0.82      3349
 macro avg      0.82      0.82      3349
 weighted avg      0.82      0.82      3349

DecisionTreeClassifier(max_depth=11)

In [17]: dt4=DecisionTreeClassifier(min_samples_leaf=10) #The minimum value of samples required to be at a leaf node.
mymodel(dt4)

precision    recall  f1-score   support

   0.00      0.82      0.83      1760
   1.00      0.81      0.80      1589

 accuracy      0.81      3349
 macro avg      0.81      0.81      3349
 weighted avg      0.81      0.81      3349

DecisionTreeClassifier(min_samples_leaf=10)

In [18]: for i in range(1,75):
    dt2=DecisionTreeClassifier(min_samples_leaf=i)
    dt2.fit(xtrain,ytrain)
    ypred=dt2.predict(xtest)
    print(f'{i}: {accuracy_score(ytest,ypred)}')

1: 0.7906837862048373
2: 0.7715736040609137
3: 0.7897879964168408
4: 0.788692206268444
5: 0.791579575928337
6: 0.7948641385488205
7: 0.8083998983687668
8: 0.8187372727983474
9: 0.8097393683487608
10: 0.8145715138847417
11: 0.8127798439817488
12: 0.8130785309047477
13: 0.8127798439817488
14: 0.8145715138847417
15: 0.8166619900567334
16: 0.82024404208719
17: 0.8214392355927143
18: 0.818751866228725
19: 0.8217378321887131
20: 0.8274311675126904
21: 0.8262167811286951
22: 0.8274311675126904
23: 0.8283009573086068
24: 0.825918184526963
25: 0.8268139743206927
26: 0.8262167811286951
27: 0.8262167811286951
28: 0.8256105879366975
29: 0.8274311675126904
30: 0.823828083607047
31: 0.8262167811286951
32: 0.8262167811286951
33: 0.8250223947447
34: 0.826516377724694
35: 0.825209913406987
36: 0.8253209913406987
37: 0.825918184526963
38: 0.8256105879366975
39: 0.828008360704688
40: 0.8262167811286951
41: 0.828001343684082
42: 0.828605539966056
43: 0.8289041504920844
44: 0.8253209913406987
45: 0.8253209913406987
46: 0.8253209913406987
47: 0.8253209913406987
48: 0.8253209913406987
49: 0.8253209913406987
50: 0.8250223947447
51: 0.8277097641086891
52: 0.8277097641086891
53: 0.824727981487031
54: 0.825204117617058
55: 0.824426201527023
56: 0.824426201527023
57: 0.823828083607047
58: 0.8214392355927143
59: 0.8214392355927143
60: 0.8214392355927143
61: 0.8214392355927143
62: 0.8214392355927143
63: 0.8214392355927143
64: 0.8214392355927143
65: 0.8217378321887131
66: 0.818604408647357
67: 0.816360934607345
68: 0.819647650917214
69: 0.818453206937201
70: 0.818453206937201
71: 0.8175974798447298
72: 0.8175974798447298
73: 0.8178560764407286
74: 0.8178560764407286
```

```
In [19]: dt5=DecisionTreeClassifier(min_samples_leaf=41)
mymodel(dt5)

precision    recall  f1-score   support

   0.00      0.86      0.81      1760
   1.00      0.80      0.85      1589

 accuracy      0.83      3349
 macro avg      0.83      0.83      3349
 weighted avg      0.83      0.83      3349

DecisionTreeClassifier(min_samples_leaf=41)

In [20]: dt6=DecisionTreeClassifier(criterion="gini",min_samples_leaf=41)
mymodel(dt6)

precision    recall  f1-score   support

   0.00      0.86      0.81      1760
   1.00      0.80      0.85      1589

 accuracy      0.83      3349
 macro avg      0.83      0.83      3349
 weighted avg      0.83      0.83      3349

DecisionTreeClassifier(min_samples_leaf=41)

In [32]: dt7=DecisionTreeClassifier(criterion="entropy",min_samples_leaf=41)
mymodel(dt7)

precision    recall  f1-score   support

   0.00      0.84      0.82      1760
   1.00      0.80      0.83      1589

 accuracy      0.82      3349
 macro avg      0.82      0.82      3349
 weighted avg      0.82      0.82      3349

DecisionTreeClassifier(criterion='entropy', min_samples_leaf=41)

In [22]: dt8=DecisionTreeClassifier(criterion="gini",max_depth=11)
mymodel(dt8)

precision    recall  f1-score   support

   0.00      0.83      0.82      1760
   1.00      0.81      0.81      1589

 accuracy      0.82      3349
 macro avg      0.82      0.82      3349
 weighted avg      0.82      0.82      3349

DecisionTreeClassifier(max_depth=11)

In [23]: dt9=DecisionTreeClassifier(criterion="entropy",max_depth=11)
mymodel(dt9)

precision    recall  f1-score   support

   0.00      0.81      0.83      1760
   1.00      0.81      0.78      1589

 accuracy      0.81      3349
 macro avg      0.81      0.81      3349
 weighted avg      0.81      0.81      3349

DecisionTreeClassifier(criterion='entropy', max_depth=11)

In [39]: from sklearn.model_selection import GridSearchCV
params = {'criterion': ['gini', 'entropy'],
          'max_depth': [3, 4, 5, 7],
          'min_samples_leaf': [10, 20, 50, 100, 150],
          }
grid_search = GridSearchCV(dt, param_grid=params)

In [41]: grid_search.fit(xtrain, ytrain)

Out[41]: GridSearchCV(estimator=DecisionTreeClassifier(),
                    param_grid={'criterion': ['gini', 'entropy'],
                                'max_depth': [3, 4, 5, 7],
                                'min_samples_leaf': [10, 20, 50, 100, 150]})

In [42]: grid_search.best_params_
Out[42]: {'criterion': 'entropy', 'max_depth': 7, 'min_samples_leaf': 50}

In [43]: my_best_preds = grid_search.predict(xtest)

In [44]: accuracy_score(ytest, my_best_preds)
Out[44]:
0.80949537152762
```

```
In [45]: print(classification_report(ytest, my_best_preds))

              precision    recall  f1-score   support

   0.00      0.81      0.83      1760
   1.00      0.81      0.78      1589

 accuracy      0.81      3349
 macro avg      0.81      0.81      3349
 weighted avg      0.81      0.81      3349

RandomForestClassifier

In [47]: from sklearn.ensemble import RandomForestClassifier
my_rf_classifier = RandomForestClassifier()

In [49]: my_rf_classifier.fit(xtrain, ytrain)

In [49]: RandomForestClassifier()

In [50]: my_predictions = my_rf_classifier.predict(xtest)

In [52]: print(accuracy_score(ytest, my_predictions))
0.841743804120633

In [53]: print(classification_report(ytest, my_predictions))

              precision    recall  f1-score   support

   0.00      0.88      0.81      1760
   1.00      0.81      0.87      1589

 accuracy      0.84      3349
 macro avg      0.84      0.84      3349
 weighted avg      0.84      0.84      3349

Voting Classifier

In [56]: from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import VotingClassifier

In [57]: df_clf = DecisionTreeClassifier()
log_clf = LogisticRegression()
svm_clf = SVC()

In [58]: voting_clf = VotingClassifier(estimators=[('lr', log_clf), ('df', df_clf), ('svc', svm_clf)])

In [60]: voting_clf.fit(xtrain, ytrain)

C:\Users\91760\anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
  https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
  https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = 17560764407286
VotingClassifier(estimators=[('lr', LogisticRegression()),
                             ('df', DecisionTreeClassifier()), ('svc', SVC())])

In [61]: from sklearn.metrics import accuracy_score

In [62]: for clf in [log_clf, df_clf, svm_clf, voting_clf]:
    clf.fit(xtrain, ytrain)
    y_pred = clf.predict(xtest)
    print(clf.__class__.__name__, accuracy_score(ytest, ypred))

C:\Users\91760\anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
  https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
  https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = 17560764407286
DecisionTreeClassifier 0.8178560764407286
SVC 0.8178560764407286
C:\Users\91760\anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
  https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
  https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = 17560764407286
VotingClassifier 0.8178560764407286
```

Bagging Classifier

```
In [64]: from sklearn.ensemble import BaggingClassifier

In [65]: bag_clf = BaggingClassifier(log_clf,n_estimators=10)

In [67]: bag_clf.fit(xtrain,ytrain)

C:\Users\91760\anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
  https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
  https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = 17560764407286
C:\Users\91760\anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
  https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
  https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = 17560764407286
BaggingClassifier(base_estimator=LogisticRegression())

In [68]: ypred = bag_clf.predict(xtest)
accuracy_score(ypred, ytest)
0.77937921769185

In [ ]:
```